

Evidence, Not Authority

Suggested responses for the worksheet. Multiple defensible answers are expected when learners state their criteria and evidence clearly.

Alicia Chua, Pawarit Laosunthara, and Eric Tang | Anyscale

The activity is not designed to force every learner to pick Take 04. It is designed to make the selection rule explicit, testable, and open to critique.

1. Blind decision

Accept any candidate supported by concrete observations. Strong responses separate visible evidence from preference. Examples include vehicle deformation, lane continuity, distant visibility, flicker, object disappearance, and unnatural plumes.

2. Operational criteria

Weather fidelity	Requested condition is recognizable, spatially coherent, and near the target severity.	Distance-dependent visibility, muted contrast, consistent particles or haze.
Structure	Road layout and important actors remain recognizable relative to the source.	Lane topology, vehicle identity, object retention, building geometry.
Stability	Appearance and motion remain coherent across frames.	Flicker, jitter, teleportation, identity breaks, transient deformation.

Good hard gates

- A vehicle, pedestrian, cone, or lane disappears or changes identity.
- Lane topology or the drivable path changes materially.
- Motion becomes physically impossible or the camera jumps.
- The requested weather is absent or obstructs all task-relevant content.
- No candidate clears a minimum acceptance threshold.

3. Lesson rubric selection

With the lesson's balanced weights of 40% weather, 35% structure, and 25% stability, Take 04 receives the highest illustrative composite score.

01	46	91	90	72.8	Stable, conservative
02	72	70	64	69.3	Atmospheric, pink cast
03	25	96	93	66.8	Sharp, weak fog
04	84	91	89	87.7	Balanced, coherent
05	50	89	87	72.9	Stable, conservative
06	58	92	90	77.9	Clean, lighter fog
07	96	54	48	69.3	Dramatic plume
08	86	68	61	73.4	Strong mood, light bloom

Important: these are fixed teaching annotations, not measured ground truth. A different defensible rubric may choose a different clip.

4. Mission profiles

- Structure-first: increase structure and stability. Take 03 or 06 may become attractive despite weaker fog.
- Weather-only: Take 07 wins because its strong plume receives no penalty for structure or instability. This is the intended proxy-failure demonstration.
- Safety-focused: apply hard gates before any average. Selecting none is valid.

What a ranking flip means

A flip is evidence that the decision depends on the stated priorities. A small perturbation causing a flip suggests low robustness. Report the runner-up and margin, not only the winner.

5. Judge the judge

Strong critiques should identify a missing criterion, a plausible exploit, and a validation plan.

- Missing properties might include object retention, lane topology, severity calibration, photometric consistency, motion realism, or downstream task usefulness.
- A candidate can score highly while hiding a critical defect if weighted averages are used without gates.
- Human calibration is needed to define fitness for purpose and review ambiguous failures.
- Deterministic tools can measure edges, depth consistency, tracking, optical flow, and flicker.
- A VLM can provide semantic evidence and failure descriptions, but can also miss subtle temporal or geometric errors.
- Abstain when every candidate fails a gate, scores fall below threshold, or evidence is too uncertain.

6. Transfer challenge

A complete response contains three observable criteria, at least one non-negotiable gate, a scoring or ranking rule, and a plan to calibrate or audit the selector.

Example: robotic manipulation

- Criteria: goal completion, object identity preservation, and physically plausible motion.
- Hard gate: the gripper passes through a solid object.
- Validation: human calibration set plus pose tracking and collision checks.

Exit ticket answers

1. N is the number of generated candidates considered by the selector.
2. Gates prevent a severe failure from being averaged away by strong scores elsewhere.
3. A selector can reward an easy proxy, miss an important property, use poorly calibrated scores, or fail to abstain.

Mastery is visible when a learner can separate generation from selection, operationalize quality, and explain why the selector itself must be evaluated.