

How Do You Test Weather You Cannot Schedule?

Alicia Chua, Pawarit Laosunthara, and Eric Tang | Anyscale

Name: _____ Date: _____

Work in this order: decide first, define quality second, calculate third, and critique the judge last.

1. Make a blind decision

Inspect the unlabeled clips before opening the scores.

Candidate I would retain: _____ Confidence (1 to 5): _____

Two visible observations that support my choice:

One failure that would make any candidate unusable:

Was your choice based on visual appeal or fitness for purpose? Explain.

2. Define fit for purpose

Write an operational definition. Name evidence that another reviewer could inspect. Assign weights that total 100.

Weather fidelity			
Structure preservation			
Temporal stability			
Optional fourth criterion			
TOTAL			100

My hard rejection gate

A candidate is rejected before scoring if:

3. Score and select

Apply hard gates first. Then compute the weighted score for eligible candidates.

score = weather weight x weather rating + structure weight x structure rating + stability weight x stability rating

01					
02					
03					
04					
05					
06					
07					
08					

Selected candidate: _____ Runner-up: _____ Margin: _____

Did this match your blind choice? Why or why not?

4. Change the mission

Create two weight profiles. Recompute or use the interactive selector.

Safety-focused evaluation					
Visual demonstration					

5. Judge the judge

What important property is missing from your rubric?

Could a candidate score highly and still be unsafe or unusable? How?

Which ratings require human calibration? Which checks could be automated?

When should the selector abstain?

6. Transfer challenge

Choose another generative task. Define three criteria, one hard gate, and one validation method.

Domain: _____

Exit ticket

1. What does N represent?
2. Why do gates come before averages?
3. How can a misaligned selector fail?